

# SOC6707 Intermediate Data Analysis

Monica Alexander

Week 8: Logistic Regression

# Notes

# Motivation

What if we are interested in modeling a binary response variable as a function of continuous and/or categorical explanatory variables?

- ▶ A binary response variable is an indicator variable that is coded 1 to indicate that an observation is a member of a particular group/category, and 0 otherwise
  - ▶ e.g. high income yes/no
  - ▶ has bachelor or higher yes/no
  - ▶ at least good self-reported health yes/no
  - ▶ life sentence yes/no
- ▶ Today we will see how we can build a regression model with a binary outcome as the dependent/response variable

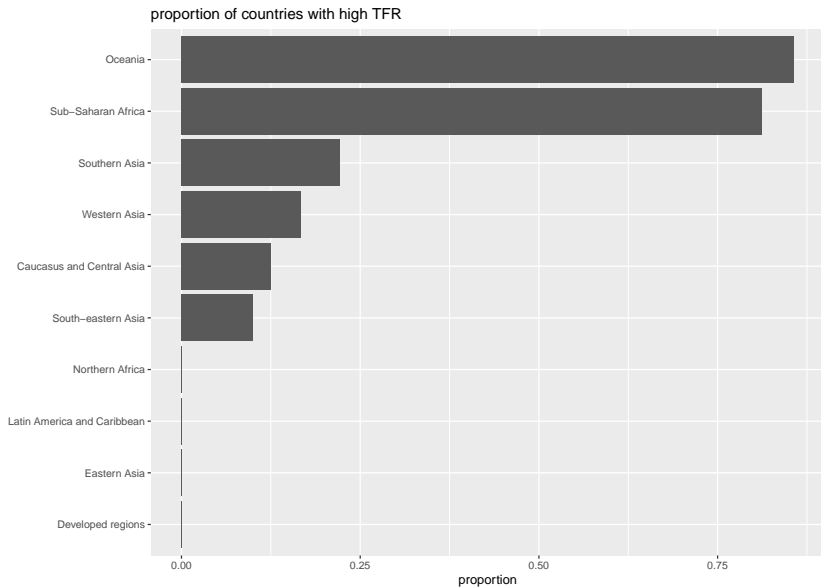
## Binary dependent variable

- ▶ For example, let's use the country indicators dataset again.
- ▶  $Y_i = 1$  if a country has a high TFR (i.e.  $TFR > 3.5$ ) and  $Y_i = 0$  otherwise.
- ▶ Note that we have to create this variable using the `ifelse` function:

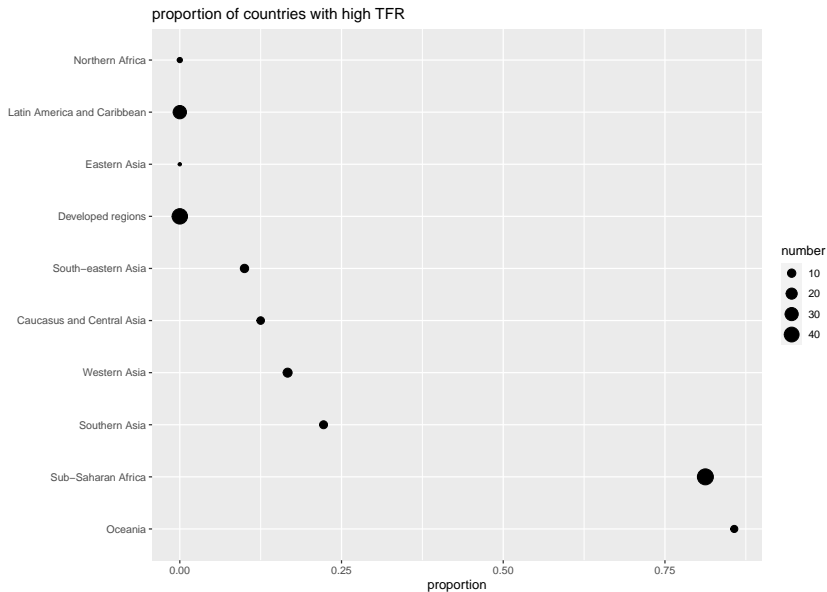
```
country_ind_2017 <- country_ind_2017 %>%  
  mutate(high_tfr = ifelse(tfr>3.5, 1, 0))  
head(country_ind_2017 %>% select(country, region, tfr, high_tfr))
```

```
## # A tibble: 6 x 4  
##   country      region      tfr high_tfr  
##   <chr>      <chr>      <dbl>   <dbl>  
## 1 Afghanistan Southern Asia  4.63     1  
## 2 Albania    Developed regions  1.64     0  
## 3 Algeria    Northern Africa  3.04     0  
## 4 Angola     Sub-Saharan Africa  5.60     1  
## 5 Antigua and Barbuda Latin America and Caribbean  2.00     0  
## 6 Argentina  Latin America and Caribbean  2.28     0
```

# Binary variables: looking at proportions is useful




## Another option



## Binary dependent variable

- ▶  $Y_i = 1$  if a country has a high TFR (i.e.  $\text{TFR} > 3.5$ ) and  $Y_i = 0$  otherwise.
- ▶ We are interested in exploring how high TFR is associated with life expectancy and gross domestic product (GDP)
- ▶ What does this actually mean, given “high TFR” is 1 or 0 (yes or no)?


???





???

high TFR versus life expectancy



## Binary dependent variable

- ▶ We are interested in exploring how high TFR is associated with life expectancy and gross domestic product (GDP)
- ▶ What does this actually mean, given “high TFR” is 1 or 0 (yes or no)?
- ▶ We are interested to see if the **probability** of high fertility is associated with life expectancy and GDP

But how do we model the probability in a regression framework?

## FANCY AVERAGES: REDUX

## The expectation of a binary variable

- ▶ Recall that the regression models we've looked at so far (SLR and MLR) are models for the **conditional expectation**  $E(Y|X)$
- ▶ Conditional expectations are just fancy averages (averages conditioning on other variables of interest)
- ▶ So if we want to model a binary outcome as a dependent variable in a regression model, we first need to find the conditional expectation

# The expectation of a binary variable

- ▶ Recall that for a discrete random variable,  $Y$ , with a known probability distribution  $P(Y_i)$  and where  $Y_i$  is the  $i$ th outcome in the set of  $k$  simple events:

$$E(Y_i) = Y_1 \times P(Y_1) + Y_2 \times P(Y_2) + \dots + Y_k \times P(Y_k) = \sum_{i=1}^k Y_i \times P(Y_i)$$

- ▶ So the expected value of a binary variable is

$$\begin{aligned} E(Y_i) &= (0)(1 - p) + (1)p \\ &= p \end{aligned}$$

- ▶ That is, the expectation of a binary variable is equal to the probability that the variable is equal to one
- ▶ So the thing that we're interested in (probability) is actually the expected value of our outcome!

# Conditional Expectation

- ▶ By extension, the conditional expectation of a binary variable is equal to the conditional probability that the variable is equal to one—that is,

$$E(Y_i | X_{i1}, \dots, X_{ik}) = P(Y_i = 1 | X_{i1}, \dots, X_{ik})$$

So we know what the conditional expectation of our outcome of interest is, so can we just do linear regression now?

## A complication

- ▶ The regression models discussed previously were direct models for the conditional expectation. But there's a complication here in that

$$E(Y_i | X_{i1}, \dots, X_{ik}) = P(Y_i = 1 | X_{i1}, \dots, X_{ik})$$

is bounded between values zero and one. It's a probability!

- ▶ We can get around this by first **transforming** the conditional expectation to be unbounded
- ▶ I.e. want to go from  $y = \text{probabilities}$  to  $y = \text{function}(\text{probabilities})$  where the function lets  $y$  be any real number.

# Logarithms

$$\log_b x$$

- ▶ The logarithm of a positive real number  $x$  with respect to base  $b$  is the exponent by which  $b$  must be raised to yield  $x$ .
- ▶ It is the inverse function to exponentiation
- ▶ The natural logarithm (often just written  $\log x$ ) is to the base  $e$ , the mathematical constant  $e \approx 2.718$

$$y = \log x$$

implies

$$x = e^y = \exp y$$

- ▶ You can think of taking the natural logarithm of  $x$  as transforming  $x$  to be on a different scale



# The logit function

- ▶ The logit function takes a probability as its argument and then returns a value between negative infinity and positive infinity
- ▶ In other words, the logit transformation of a probability is unbounded even though the probability is bounded by the unit interval,  $[0,1]$
- ▶ It is also called log-odds


The logit function of probability  $p$  is

$$\text{logit } p = \log \frac{p}{1-p}$$

# The logit function

For example,

$$\text{logit } 0.5 = \log \frac{0.5}{1 - 0.5} = \log 1 = 0$$



## Aside: odds

Given probability  $p$ , odds are calculated as

$$\frac{p}{1 - p}$$

- ▶ Odds provide a measure of the likelihood of a particular outcome. They are calculated as the ratio of the number of events that produce the outcome to the number that don't.
- ▶ Another way of expressing likelihood
- ▶ Often expressed as "1 to x"
- ▶ e.g. six sided die:
  - ▶ Probability rolling a 6 = ?
  - ▶ Odds of rolling a 6 = ?

## Back to our problem

We want to model the conditional expectation

$$E(Y_i \mid X_{i1}, \dots, X_{ik}) = P(Y_i = 1 \mid X_{i1}, \dots, X_{ik})$$

For simplicity/ease of reading, I'm going to use

$$p = P(Y_i = 1 \mid X_{i1}, \dots, X_{ik})$$

# The logistic regression model

Logistic regression is a model for the conditional expectation of a binary response variable—that is, for the conditional probability that a binary response variable is equal to one.

$$\log \left( \frac{p}{1-p} \right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik}$$

where  $\log \left( \frac{p}{1-p} \right)$  is known as the “log odds,” or the “logit” transformation, and the  $\beta$  are unknown parameters to be estimated from data

# The logistic regression model

$$\log \left( \frac{p}{1-p} \right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik}$$

We can rearrange this formula to get an expression for probability that  $Y_i = 1$ :

$$p = \frac{\exp(\beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik})}{1 + \exp(\beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik})}$$

- ▶ This is the inverse of the logit function
- ▶ The inverse of the logit link function is bounded by the unit interval (i.e., it falls between 0 and 1 for any value), which ensures that the conditional probabilities all fall within the logical range

# The logistic regression model

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik}$$

To summarize:

- ▶ We transform probabilities to run a regression model that can have values anywhere on the real line
- ▶ We can then untransform these probabilities to get values back on the  $[0,1]$  scale

## Interpreting logistic regression on the logit scale

$$\log \left( \frac{p}{1-p} \right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_k X_{ik}$$

What is  $\beta_0$ ? Now I have to go back to the full notation:

$$\log \left( \frac{P(Y_i=1|X_{i1}=0, \dots, X_{ik}=0)}{1-P(Y_i=1|X_{i1}=0, \dots, X_{ik}=0)} \right) = \beta_0 + \beta_1(0) + \cdots + \beta_k(0) \\ = \beta_0$$

$\beta_0$  is the log odds that  $Y_i = 1$  given that all explanatory variables are equal to zero.



# Interpreting logistic regression on the logit scale

What is  $\beta_1$ ?

$$\begin{aligned} & \log \left( \frac{P(Y_i = 1 \mid X_{i1} = x_1^* + 1, X_{i2} = x_2^*, \dots, X_{ik} = x_k^*)}{1 - P(Y_i = 1 \mid X_{i1} = x_1^* + 1, X_{i2} = x_2^*, \dots, X_{ik} = x_k^*)} \right) \\ & - \log \left( \frac{P(Y_i = 1 \mid X_{i1} = x_1^*, X_{i2} = x_2^*, \dots, X_{ik} = x_k^*)}{1 - P(Y_i = 1 \mid X_{i1} = x_1^*, X_{i2} = x_2^*, \dots, X_{ik} = x_k^*)} \right) \\ & = (\beta_0 + \beta_1(x_1^* + 1) + \beta_2 x_2^* + \dots + \beta_k x_k^*) - (\beta_0 + \beta_1 x_1^* + \beta_2 x_2^* + \dots + \beta_k x_k^*) \\ & = \beta_1 \end{aligned}$$

$\beta_1$  is a log odds ratio, which gives the change in the log odds that  $Y_i = 1$  associated with a unit increase in  $X_{i1}$ , holding other variables constant

## Interpreting logistic regression on the odds scale

$$\log \left( \frac{P(Y_i = 1 \mid X_{i1}, \dots, X_{ik})}{1 - P(Y_i = 1 \mid X_{i1}, \dots, X_{ik})} \right) = \beta_0 + \beta_1 X_{i1} + \dots + \beta_k X_{ik}$$

What is  $\exp \beta_0$ ?

$$\begin{aligned} \exp \beta_0 &= \exp \left( \log \left( \frac{P(Y_i=1|X_{i1}=0,\dots,X_{ik}=0)}{1-P(Y_i=1|X_{i1}=0,\dots,X_{ik}=0)} \right) \right) \\ &= \frac{P(Y_i=1|X_{i1}=0,\dots,X_{ik}=0)}{1-P(Y_i=1|X_{i1}=0,\dots,X_{ik}=0)} \end{aligned}$$

$\exp \beta_0$  is the odds that  $Y_i = 1$  given that all explanatory variables are equal to zero.

## Interpreting logistic regression on the odds scale

What is  $\exp \beta_1$ ?

$$\exp(\beta_1) = \frac{P(Y_i = 1 \mid X_{i1} = x_1^* + 1, \dots)}{1 - P(Y_i = 1 \mid X_{i1} = x_1^* + 1, \dots)} / \frac{P(Y_i = 1 \mid X_{i1} = x_1^*, \dots)}{1 - P(Y_i = 1 \mid X_{i1} = x_1^*, \dots)}$$

$\exp \beta_1$  is a odds ratio, which ratio of the odds that  $Y_i = 1$  associated with a unit increase in  $X_{i1}$ , holding other variables constant

## Example in R

- ▶ Can run logistic regression in R using the `glm` function
- ▶ The additional `family` argument is related to the fact we are dealing with a binary response variable

```
lr_mod <- glm(high_tfr ~ life_expectancy + gdp,  
              family = "binomial", data = country_ind_2017)
```

# Example in R

```
summary(lr_mod)
```

```
##
## Call:
## glm(formula = high_tfr ~ life_expectancy + gdp, family = "binomial",
##      data = country_ind_2017)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.08570  -0.23518  -0.02127   0.18777   2.36080
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   21.9815193   4.4298761    4.962 6.97e-07 ***
## life_expectancy -0.2960674   0.0627786   -4.716 2.40e-06 ***
## gdp           -0.0002081   0.0000654   -3.181 0.00147 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 211.886  on 175  degrees of freedom
## Residual deviance:  73.653  on 173  degrees of freedom
## AIC: 79.653
##
## Number of Fisher Scoring iterations: 8
```

# Questions

## Interpret

▶  $\beta_1$

▶  $\exp(\beta_1)$

```
coef(lr_mod)
```

```
##      (Intercept) life_expectancy      gdp
## 21.9815193435   -0.2960674332   -0.0002080662
```

```
exp(coef(lr_mod))
```

```
##      (Intercept) life_expectancy      gdp
## 3.519270e+09    7.437373e-01    9.997920e-01
```

# Questions

- What is the probability of high TFR for a country with a life expectancy of 70 and a GDP of 9500?

```
beta0 <- coef(lr_mod)[[1]] # used double square brackets here to remove names (could use single)  
beta1 <- coef(lr_mod)[[2]]  
beta2 <- coef(lr_mod)[[3]]
```

```
estimated_log_odds <- beta0 + beta1*70 + beta2*9500
```

```
estimated_probability <- exp(estimated_log_odds)/(1+exp(estimated_log_odds))
```

```
estimated_log_odds
```

```
## [1] -0.7198301
```

```
estimated_probability
```

```
## [1] 0.3274304
```

# Including a categorical explanatory variable

```
lr_mod_2 <- glm(high_tfr ~ region,
                 family = "binomial", data = country_ind_2017)
summary(lr_mod_2)
```

```
##
## Call:
## glm(formula = high_tfr ~ region, family = "binomial", data = country_ind_2017)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.97277  -0.00005  -0.00005   0.64442   2.14597
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -1.9459     1.0690  -1.820  0.06872 .
## regionDeveloped regions      -18.6202    2672.9540  -0.007  0.99444
## regionEastern Asia      -18.6202    10236.6339  -0.002  0.99855
## regionLatin America and Caribbean      -18.6202    3184.4686  -0.006  0.99533
## regionNorthern Africa      -18.6202     8865.1850  -0.002  0.99832
## regionOceania           3.7377     1.5197   2.459  0.01391 *
## regionSouth-eastern Asia      -0.2513     1.5013  -0.167  0.86706
## regionSouthern Asia          0.6931     1.3363   0.519  0.60397
## regionSub-Saharan Africa       3.4122     1.1312   3.016  0.00256 **
## regionWestern Asia          0.3365     1.3202   0.255  0.79882
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 211.886  on 175  degrees of freedom
## Residual deviance:  84.947  on 166  degrees of freedom
## AIC: 104.95
##
## Number of Fisher Scoring iterations: 19
```



# Categorical explanatory variables

- ▶ The coefficient on Sub-Saharan Africa is 3.41. What does this mean?

```
exp(coef(lr_mod_2)[9])
```

```
## regionSub-Saharan Africa  
##                30.33333
```

# Inference

## Some brief comments on the sampling distribution estimators $\hat{\beta}$

- ▶ Thinking back to SLR and MLR, if we believed the 5 assumptions stated, we could write down the sampling distribution for the estimator  $\hat{\beta}$
- ▶ (It was a Normal distribution)
- ▶ We could then use this property to make inferences about how likely  $\hat{\beta}$  was to be different from zero, for example (hypothesis testing)
- ▶ We can use a similar approach here with logistic regression

## Asymptotic distribution of $\hat{\beta}$

- ▶ It is known that the limiting distribution of  $\hat{\beta}_k$  is normal with a mean  $\beta_k$  and some variance (related to the properties of the estimator)
- ▶ Because the probability distribution of  $\hat{\beta}_k$  converges to a normal distribution as the sample size increases, we can use this fact to make approximate inferences about  $\beta_k$
- ▶ It turns out that the standardized version

$$Z_{\hat{\beta}_k} = \frac{\hat{\beta}_k - \beta_k}{se(\hat{\beta}_k)}$$

follows a standard normal distribution, which we can use to make inferences about  $\beta_k$

# Hypothesis testing

- ▶ The  $\beta_k$  parameters are unknown population quantities of interest, which we have estimated with data from a random sample of the population
- ▶ We can test hypotheses about these unknown population quantities based on the fact that their standardized estimates follow an approximately standard normal distribution in large samples
- ▶ With knowledge of the distribution of  $Z_{\hat{\beta}_k}$  we can make probabilistic statements about the chances of observing any particular value of  $Z_{\hat{\beta}_k}$  given a hypothesized value for the unknown parameter of interest
- ▶ As before, we are usually testing the null hypothesis that  $\beta_k = 0$
- ▶ This test is called the Wald test

# The Wald test

1. State your null and alternative hypotheses about  $\beta_k$
2. Choose the level of type-I error,  $\alpha$
3. Compute the Wald test statistic  $z_{\hat{\beta}_k} = \frac{\hat{\beta}_k - \beta_k}{se(\hat{\beta}_k)}$
4. Compute the p-value, which gives the probability of observing a test statistic as or more extreme than  $z_{\hat{\beta}_k}$  under the assumption that the null hypothesis is true
5. Make a decision (reject the null if the p-value is less than  $\alpha$ , and fail to reject otherwise)

Reminder: think of the p-value as a summary measure of 'evidence against the null hypothesis' (and *not* as evidence for the alternative hypothesis)

# Example

```
summary(lr_mod)
```

```
##
## Call:
## glm(formula = high_tfr ~ life_expectancy + gdp, family = "binomial",
##      data = country_ind_2017)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.08570  -0.23518  -0.02127   0.18777   2.36080
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   21.9815193   4.4298761    4.962 6.97e-07 ***
## life_expectancy -0.2960674   0.0627786   -4.716 2.40e-06 ***
## gdp            -0.0002081   0.0000654   -3.181 0.00147 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 211.886  on 175  degrees of freedom
## Residual deviance:  73.653  on 173  degrees of freedom
## AIC: 79.653
##
## Number of Fisher Scoring iterations: 8
```

# Summary

- ▶ Logistic regression can be used when the outcome of interest is binary (yes/no)
  - ▶ Aside: why logistic?
- ▶ You can have one or more explanatory variables, which can be quantitative or categorical
- ▶ Practically, running logistic regression in R is very similar to linear regression
- ▶ Interpretation is usually a bit harder